
Vision-Language Pretraining for Image-to-Text Generation

Pranshul Bhatnagar, Aesha Gandhi, Gaurav Law //

Abstract

Recent advances in Vision-Language Pretraining (VLP) have greatly improved bidirectional understanding between visual and textual modalities. Models such as Bootstrapped Language Image Pretraining (BLIP) and Vision Transformer (ViT) now excel at tasks such as visual question answering (VQA) and image retrieval. However, while text-to-image generation has received significant attention, the reverse task of generating rich, context-aware text from images remains less explored and more challenging. This project investigates whether VLP models can effectively invert this relationship by producing accurate, semantically meaningful descriptions from images. We focus on ViT and BLIP due to their central role in current text-to-image pipelines, training them on approximately 10,000 image-prompt pairs from DiffusionDB. Performance is assessed using established metrics such as BLEU and BERTScore to measure linguistic quality and semantic alignment.

1 Introduction

Recent advances in Vision-Language Pretraining (VLP) have greatly improved bidirectional understanding between visual and textual modalities. Models such as Bootstrapped Language Image Pre-training (BLIP) and Vision Transformer (ViT) now excel at tasks such as visual question answering (VQA) and image retrieval. However, while text-to-image generation has received significant attention, the reverse task — generating rich, context-aware text from images — remains less explored and more challenging.

This project investigates whether VLP models can effectively invert this relationship by producing accurate, semantically meaningful descriptions from images. We focus on ViT and BLIP due to their central role in current text-to-image pipelines. ViT serves as a pure vision encoder, whereas BLIP integrates both image and text inputs within a multimodal training loop. BLIP's architecture wraps ViT within a framework that learns joint representations, making it well-suited for our reverse-generation task.

Using pre-collected image–prompt pairs, we train models to generate new captions, then evaluate them through validation and testing. Performance is assessed using established metrics such as BLEU and BERTScore to measure linguistic quality and semantic alignment.

Our [Github Repo](#) contains all the code for this project along with the finetuned models.

2 Data

For the purpose of this project, we sourced our data from DiffusionDB, a large-scale text-to-image dataset that contains over 14 million images generated with Stable Diffusion 2.0 using user-provided prompts and hyperparameters. From this collection, we randomly selected approximately 10,000 images, each with a resolution of 512×512 pixels, along with their corresponding English prompts, which varied widely in length (distribution shown in Figure 1). Random sampling allowed us to

37 capture a broad and unbiased subset of the dataset while keeping the computational load manageable.
 38 Although more targeted sampling strategies could have revealed specific trends, using a random
 39 selection helped ensure that our results remained generalizable across a wide range of prompt-image
 40 pairs.

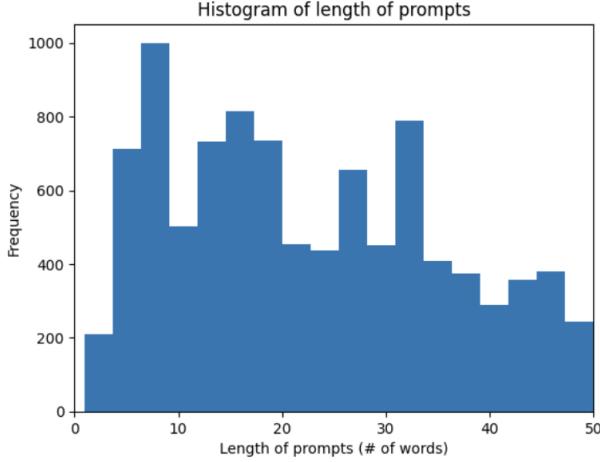


Figure 1: Distribution of Length of prompts in our Dataset

41 Following this, the dataset was split into train (60%), validation (20%), and test (20%) sets using
 42 a fixed seed of 42 to ensure reproducibility. The training set was used to teach the models to learn
 43 patterns between images and their associated prompts. The validation set was used to monitor the
 44 model’s performance during training, allowing us to fine tune hyperparameters and prevent overfitting.
 45 Finally, the testing set was used to evaluate the model’s performance on unseen data, ensuring an
 46 unbiased assessment of generalization.

47 3 Model Architecture

48 3.1 Vision Transformer + GPT-2

49 The Vision Transformer applies the transformer architecture by splitting images into small patches
 50 (such as 16×16 pixels), flattening them, and projecting into embeddings so each patch becomes
 51 a “token.” Positional embeddings preserve spatial structure, and the sequence of patch tokens is
 52 processed by a transformer encoder (12 transformer blocks) whose self-attention layers build a
 53 global understanding of the image. The resulting visual embedding is fed into the GPT-2 decoder
 54 (12 transformer blocks with cross-attention), which generates the caption one word at a time as an
 55 autoregressive model, forming the encoder–decoder setup. Inference is performed using beam search
 56 with four beams and early stopping.

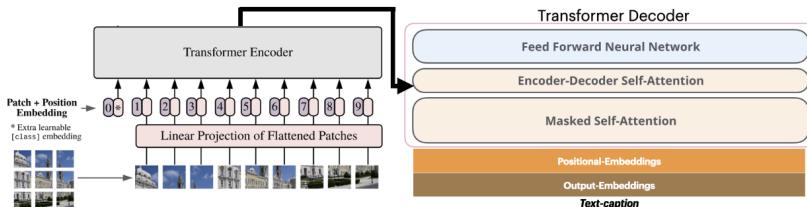


Figure 2: ViT-GPT2 Model Architecture

57 3.2 Bootstrapping Language-Image Pre-Training

58 Bootstrapping Language-Image Pre-Training is a generative method using an encoder-decoder
 59 architecture enhanced by a caption-filtering mechanism (Li et al., 2022). BLIP is designed to

60 generate accurate image captions while filtering out noisy image-text pairs, allowing it to effectively
 61 leverage large-scale pre-training data. Its Multimodal Encoder-Decoder (MED) model consists of:
 62 (1) Unimodal encoder trained with image-text contrastive (ITC) loss to align visual and textual
 63 representations, (2) Image-grounded text encoder which incorporates visual information through
 64 cross-attention layers and is trained with image-text matching (ITM) loss to distinguish correct from
 65 incorrect image-text pairs, and (3) Image-grounded text decoder that generates captions using a
 66 language modeling (LM) loss, employing causal self-attention along with cross-attention layers to
 67 capture interactions between vision and language.

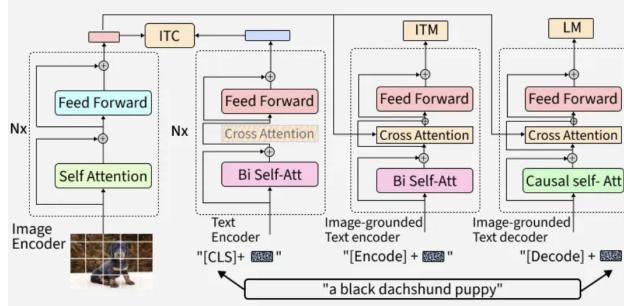


Fig 3: BLIP Model Architecture

Figure 3: BLIP Model Architecture

68 BLIP has shown strong performance in both vision-language understanding and generation tasks;
 69 because of its robust pre-training approach and proven results, we selected BLIP as one of the
 70 candidates for our experimentation.

71 4 Methodology

72 Modern image-to-text systems typically rely on neural encoder-decoder architectures, such as
 73 transformer-based models, which excel at generating coherent text sequences from visual inputs.
 74 Because our images were produced through a similar generative process, leveraging transformers
 75 for caption generation offers a natural alignment between the image representations and the text
 76 decoding pipeline.

77 To strengthen the model’s performance, we incorporated both pre-training — enabling the model to
 78 learn broad visual and linguistic patterns — and transfer learning, allowing us to adapt those learned
 79 features to our specific dataset with far fewer resources. Together, these techniques improve accuracy,
 80 reduce training time, and increase overall effectiveness in text generation.

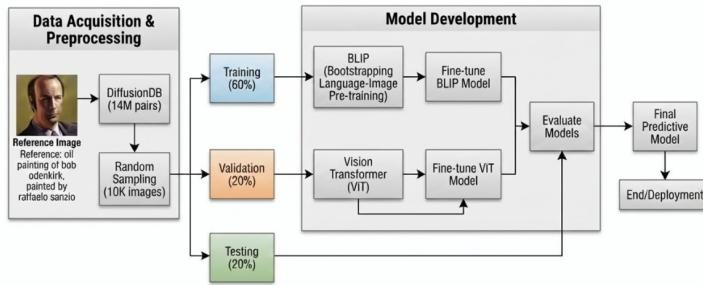


Figure 4: Flowchart outlining our Methodology

81 We implemented a PyTorch Dataset class to handle loading and preprocessing, converting images into
 82 tensors and captions into tokenized sequences using the appropriate processor. The dataset was split
 83 into training and validation sets, shuffled to ensure randomness, and loaded in batches to optimize
 84 GPU memory usage. Training for both models was performed on a NVIDIA A100 GPU using a
 85 sequence-to-sequence objective with cross-entropy loss.

86 **5 Evaluation Metric**

87 To evaluate the performance of both the BLIP-and-ViT-based captioning models, we developed a
88 custom metric that combines Bilingual Evaluation Understudy (BLEU) and Bidirectional Encoder
89 Representations (BERT) Score to capture semantic meaning and lexical accuracy. BLEU measures
90 n-gram overlap between the generated caption and the reference caption, which directly evaluates
91 the exactness of the output. BERT Score computes similarity using contextual embeddings from a
92 transformer encoder, allowing it to capture deeper semantic alignment even when the wording differs.
93 To emphasize semantic correctness without ignoring lexical precision, we assigned 65% weight to
94 BERTScore and 35% weight to BLEU:

$$Score = 0.65 \times BERTScore + 0.35 \times BLEU \quad (1)$$

95 All metrics were computed on cleaned predictions for consistency. Both the pretrained baseline and
96 the fine-tuned model used the same inference pipeline.

97 **6 Results & Discussions**

98 **6.1 BLIP**

99 For fine-tuning BLIP, we used the pre-trained model Salesforce/blip-image-captioning-base and
100 trained it for 10 epochs. Fine-tuning allowed the model to adjust its pre-learned visual and linguistic
101 representations to better align with the specific distribution and style of our dataset. We employed a
102 learning rate of 1e-4 with a CosineAnnealingWarmRestarts scheduler, which periodically increases
103 the learning rate to help the model escape local minima. A weight decay of 1e-6 was used, and
104 gradient accumulation over 4 steps allowed for effective batch size scaling. The maximum sequence
105 length for captions was set to 128 tokens.

106 During fine-tuning, both training and validation losses decreased steadily, reflecting the model’s
107 ability to learn from the dataset while generalizing well to unseen samples. As shown in Figure 5a,
108 the validation loss consistently followed the trend of the training loss, and the lowest validation
109 loss was achieved at the final epoch, indicating stable and effective fine-tuning without evidence of
110 overfitting. The fine-tuned BLIP model achieved a lowest validation loss of 0.0075 by the end of the
111 10th epoch.

(a) BLIP Fine-tuning Results

Metric	Pre	Fine	Change	Improv
BERT	0.44	0.76	+0.32	+73%
BLEU	0.27	0.78	+0.51	+189%
Comb.	0.38	0.77	+0.39	+102%

(b) ViT GPT2 Fine-tuning Results

Metric	Pre	Fine	Change	Improv
BERT	0.42	0.43	+0.01	+3%
BLEU	0.00	0.00	+0.00	+163%
Comb.	0.27	0.28	+0.01	+3%

Table 1: Comparison of fine-tuning results for both models

112 The fine-tuned BLIP model shows substantial improvements over the pretrained baseline as shown
113 in Table 1a. BERTScore increased by over 73%, indicating that the fine-tuned model produced
114 captions that were much more contextually aligned with the reference text. BLEU improved the
115 most, rising by nearly 189%, showing that the model learned to generate more precise and textually
116 accurate descriptions. When combined, the fine-tuned model achieved an improvement of about
117 102%, demonstrating that fine-tuning significantly enhanced both the semantic and lexical quality of
118 the generated captions.

119 When we looked at BLIP’s predictions manually (Table 2), there’s a clear difference between the
120 captions it handled well and the ones it struggled with. For the good predictions, the model did a
121 strong job capturing both the overall theme of the image and small details, especially when the image
122 had a clear subject, a distinct art style, or strong visual cues. The bad predictions show some of
123 BLIP’s weaknesses; in images that were more ambiguous or had fewer visual anchors, the model
124 tended to guess or hallucinate details that were not actually present.

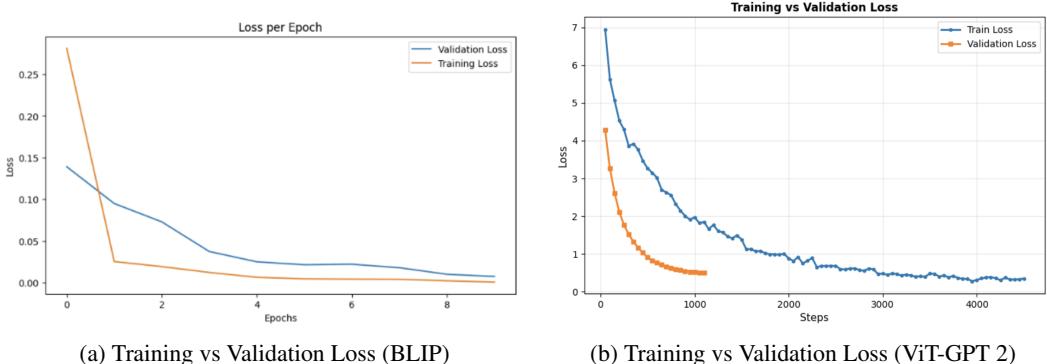


Figure 5: Training and validation loss curves for both models

Table 2: Example Prompt Predictions using our Fine-tuned BLIP Model

Image	Prompt	Caption	BLEU	BERT	Score
GOOD PREDICTIONS					
	A portrait of a beautiful female cyborg with a cracked porcelain face by wlop, exposed inner structure, brightly glowing white eyes, art nouveau card, trending on artstation	a portrait of a beautiful female cyborg wearing an intricate venetian mask by Wlop	0.79	0.69	0.73
	frontal portrait of a muscular woman in a long - sleeved linen shirt, wearing vintage engineer goggles, by norman rockwell	a portrait of a muscular anthropomorphic cyberpunk iguana! in leather spacesuit	0.51	0.58	0.52
	top-heavy 20 year old with messy black hair and big beard eats mayonnaise straight out of the jar with his bare hands	obese man with messy black hair and big beard drinking mayonnaise from the jar	0.64	0.76	0.72
BAD PREDICTIONS					
	jack vettriano painting of lilly collins	a beautiful british woman with short brown hair, gentle, vibrant amber eyes, standing on a rooftop	0.17	0.22	0.1999
	A house with a beautiful backyard.	photorealistic beautiful cherry blossom forest with victorian lanterns lining the stone pathway. hyperdetailed photo	0.01	0.51	0.33

125 6.2 ViT

126 Our dataset images were converted to RGB and processed using the ViTImageProcessor, which
 127 handled normalization and resizing to the pretrained ViT encoder’s expected input size. Captions
 128 were tokenized using the GPT-2 tokenizer with a maximum length of 80 tokens; padding tokens were
 129 replaced with -100 placeholders to exclude them from loss computation.

130 Fine-tuning used the pretrained nlpconnect/vit-gpt2-image-captioning VisionEncoderDecoderModel.
 131 Training was conducted for six epochs with a per-device batch size of four and gradient accumulation
 132 over two steps for an effective batch size of eight. The AdamW optimizer was used with a learning
 133 rate of 5×10^{-5} , and mixed-precision FP16 training was enabled for memory efficiency on GPU.

Table 3: Example Prompt Predictions using our Fine-tuned ViT-GPT2 Model

Image	Prompt	Caption	BLEU	BERT	Score
GOOD PREDICTIONS					
	a close up photo of a glass duck on a table, visible background, professional photography	close photo a duck a of glass on table visible, background professional	0.43	0.72	0.48
	photo still of al roker in a ball pit!!!!!! at age 4 6 years old 4 6 years of age!!!!!!	still christ a loft !!!!!!! age 5 old fun pop by bag bund !!!!!!! age 5 old 5 of!!!!!!	0.09	0.64	0.44
	A portrait of a beautiful anime girl wearing a long coat, shoulder length straight black hair	cute girl a anime wearing stylish coat shoulder straight, length bob hair straight hair	0.01	0.66	0.43
BAD PREDICTIONS					
	Alstom tr train on track	blue lights	0.00	0.20	0.13
	amazing futuristic design for a 3D printer, intricate, highly detailed, smooth, sharp focus	of all	0.01	0.20	0.13

134 As shown in Figure 5b, training loss declined smoothly, from ~6.9 at initialization to below 1.0
 135 by epoch 6, and validation loss followed the same downward trend, indicating successful learning
 136 and stable generalization. As shown in Table 1b, the fine-tuned ViT captioning model shows small
 137 improvements over the pretrained baseline. While the absolute values remain modest, the consistent
 138 upward movement across metrics shows that the ViT model did benefit from task-specific training
 139 and became better aligned with the captioning style of the dataset. However, a qualitative review of
 140 the example outputs (Table 3) shows that the model still struggles with generating detailed and fully
 141 accurate descriptions, occasionally generating repetitive or incomplete phrases.

142 We also noticed that images with higher BERTScores tended to produce captions that were more
 143 coherent and better aligned with the intent of the prompt, even in cases where the combined score
 144 was more moderate. This suggests that BERTScore captures semantic alignment particularly well
 145 in our setting and better reflected whether a caption made sense to a human reader, even when the
 146 wording differed from the ground truth.

147 7 Conclusion

148 This project investigated whether the text-to-image process can be meaningfully reversed by gen-
 149 erating detailed, context-aware descriptions from images. Comparing a ViT-GPT encoder-decoder
 150 model with the multimodally pre-trained BLIP model revealed a clear advantage for models grounded
 151 in joint vision-language pre-training. BLIP showed substantial gains after fine-tuning, capturing
 152 objects, styles, and visual details more effectively, whereas the ViT-GPT model achieved only modest
 153 improvements due to its limited semantic grounding.

154 Despite BLIP’s strong performance, both models struggled with deeper contextual reasoning, abstract
 155 concepts, and complex relationships, often producing captions shorter and less expressive than target
 156 prompts. These limitations highlight the need for richer training data, improved architectures, and
 157 more semantically sensitive evaluation metrics.

158 Overall, our findings demonstrate that reversing the text-to-image relationship is achievable but highly
 159 dependent on the model’s pre-training paradigm. Advancing this direction through larger datasets,
 160 better tokenization, refinement techniques, and architectures built for bidirectional vision language

161 tasks holds promising potential for applications in accessibility, robotics, retrieval, and multimodal
162 AI systems.

163 **References**

- 164 [1] A. Kumar, "The Illustrated Image Captioning using transformers," *Ankur | NLP Enthusiast*, Nov. 20, 2022.
165 <https://ankur3107.github.io/blogs/the-illustrated-image-captioning-using-transformers/> (accessed Mar. 08, 2023)
- 166 [2] Li, J., Li, D., Xiong, C., & Hoi, S. (2022). Blip: Bootstrapping language-image pre-training for unified vision-
167 language understanding and generation. In *International Conference on Machine Learning* (pp. 12888-12900).
168 PMLR.
- 169 [3] Ordonez, V., Kulkarni, G., & Berg, T. (2011). Im2text: Describing images using 1 million captioned
170 photographs. *Advances in neural information processing systems*, 24.
- 171 [4] Radford A, Kim JW, Hallacy C, et al. Learning transferable visual models from Natural Language Supervision.
172 arXiv.org. <https://arxiv.org/abs/2103.00020>. Published February 26, 2021. Retrieved April 5, 2023.
- 173 [5] Reimers N, Gurevych I. Sentence-bert: Sentence embeddings using Siamese Bert-Networks. arXiv.org.
174 <https://arxiv.org/abs/1908.10084>. Published August 27, 2019. Accessed April 12, 2023.
- 175 [6] Singhal, A. (2001). Modern information retrieval: A brief overview. *IEEE Data Eng. Bull.*, 24(4), 35-43.
- 176 [7] Wikimedia Foundation. (2023, February 26). Cosine similarity. *Wikipedia*. Retrieved April 5, 2023, from
177 https://en.wikipedia.org/wiki/Cosine_similarity
- 178 [8] Zhu Y, Lu S, Zheng L, et al. Texxygen: A benchmarking platform for text generation models. arXiv.org.
179 <https://arxiv.org/abs/1802.01886>. Published February 6, 2018. Accessed April 15, 2023.